

M.Sc. (Five Year Integrated) in Computer Science (Artificial Intelligence & Data Science)

Seventh Semester

Laboratory Record

21-805-0704: Computational Linguistics Lab

*Submitted in partial fulfillment
of the requirements for the award of degree in
Master of Science (Five Year Integrated)
in Computer Science (Artificial Intelligence & Data Science) of
Cochin University of Science and Technology (CUSAT)
Kochi*



Submitted by

ASHITHA P
(80521006)

DEPARTMENT OF COMPUTER SCIENCE
COCHIN UNIVERSITY OF SCIENCE AND TECHNOLOGY (CUSAT)
KOCHI-682022

December 2024

DEPARTMENT OF COMPUTER SCIENCE
COCHIN UNIVERSITY OF SCIENCE AND TECHNOLOGY (CUSAT)
KOCHI, KERALA-682022



*This is to certify that the software laboratory record for **21-805-0704: Computational Linguistics Lab** is a record of work carried out by **ASHITHAP(80521006)**, in partial fulfillment of the requirements for the award of degree in **Master of Science (Five Year Integrated) in Computer Science (Artificial Intelligence & Data Science)** of Cochin University of Science and Technology (CUSAT), Kochi. The lab record has been approved as it satisfies the academic requirements in respect of the second semester laboratory prescribed for the Master of Science (Five Year Integrated) in Computer Science degree.*

Faculty Member in-charge

Dr. Jeena K
Assistant Professor
Department of Computer Science
CUSAT

Dr. Madhu S. Nair
Professor and Head
Department of Computer Science
CUSAT

Table of Contents

Sl.No.	Program	Pg.No.
1	Text Tokenizer	pg 1
2	Finite State Automata	pg 3
3	Finite State Transducer	pg 6
4	Minimum Edit Distance	pg 9
5	Spell Checker	pg 11
6	Sentiment Analysis	pg 15
7	POS Tagging	pg 18
8	Bigram Probability	pg 21
9	TF-IDF Matrix	pg 23
10	PPMI Matrix	pg 25
11	Naive Bayes Classifier	pg 28

Text Tokenizer

AIM

Implement a simple rule-based Text tokenizer for the English language using regular expressions. Your tokenizer should consider punctuations and special symbols as separate tokens. Contractions like "isn't" should be regarded as 2 tokens - "is" and "n't". Also identify abbreviations (eg, U.S.A) and internal hyphenation (eg. ice-cream), as single tokens.

PROGRAM

```
def tokenizeText(text):
    pattern = r"""
    (?:[A-Za-z]\.){2,}[A-Z]   # Matches abbreviations like U.S.A or U.K
    | \w+(?:-\w+)+            # Matches internal hyphenation like ice-cream
    | \b([A-Za-z]+)(n't|'s|'ll|'em|'ve|'re|'d)\b # Matches contractions like isn't,
you'll
    | \b\w+\b                # Matches standalone words
    | [.,!?"()\\[\]{}<>]    # Matches punctuation as separate tokens
    """
    tokens = []
    for match in re.finditer(pattern, text, flags=re.VERBOSE):
        if match.group(1):
            tokens.extend([match.group(1), match.group(2)])
        else:
            tokens.append(match.group(0))
    return tokens

text = 'Implement a simple rule-based Text tokenizer for the English language using
regular expressions. Your tokenizer should consider punctuations and special
symbols as separate tokens. Contractions like "isn\'t" should be regarded as 2 tokens
- "is" and "n\'t". Also identify abbreviations (eg, U.S.A) and internal hyphenation
(eg. ice-cream), as single tokens.'
tokens = tokenizeText(text)
tokens = set(tokens)
print(tokens)
```

SAMPLE INPUT-OUTPUT

```
{'.', 'regular', 'symbols', 'simple', ')', 'separate', 'Text', ',,', 'as',  
  'rule-based', 'abbreviations', 'hyphenation', 'for', 'special', 'tokenizer',  
  'U.S.A', 'punctuations', 'identify', 'consider', "n't", 'single', 'regarded',  
  'ice-cream', 'the', 'and', 'Contractions', 'Implement', 'expressions',  
  'using', '2', 'e.g.', 'a', 'be', 't', 'tokens', 'should', '"', 'language',  
  'n', 'like', 'is', 'Also', '(', 'English', 'Your', 'internal'}
```

Finite State Automata

AIM

Design and implement a Finite State Automata(FSA) that accepts English plural nouns ending with the character 'y', e.g. boys, toys, ponies, skies, and puppies but not boies or toies or ponys. (Hint: Words that end with a vowel followed by 'y' are appended with 's' and will not be replaced with "ies" in their plural form).

PROGRAM

```
import graphviz

def is_plural_noun_accepted_fsa(word):
    # Check if the word is at least 2 characters long and ends with 's'
    if len(word) < 2 or word[-1] != 's':
        return False

    # Reverse the word for easier state processing
    word = word[::-1]
    state = 'S1' # Initial state

    for char in word[1:]: # Skip the last character ('s')
        if state == 'S1':
            if char == 'y':
                state = 'S2' # Transition to state for vowel + 'y'
            elif char == 'e':
                state = 'S3' # Transition to state for 'ies'
            else:
                return False # Invalid transition

        elif state == 'S2':
            if char in 'aeiou':
                state = 'S5' # Valid transition to vowel + 'y' state
            else:
                return False # Invalid transition

        elif state == 'S3':
            if char == 'i':
                state = 'S4' # Valid transition to state for 'ies'
            else:
                return False # Invalid transition
```

```
        elif state == 'S4':
            if char.isalpha() and char not in 'aeiou':
                state = 'S6' # Transition to final state after consonant
            else:
                return False # Invalid transition

        # States S5 and S6 indicate acceptance; continue checking
        elif state in ('S5', 'S6'):
            continue

    return True # If we reach here, the word is a valid plural noun

def draw_fsa():
    dot = graphviz.Digraph(comment='FSA for Plural Noun Detection')

    # Define states
    dot.node('S1', 'S1 (Start)', shape='circle') # Start state
    dot.node('S2', 'S2 (y)', shape='circle')
    dot.node('S3', 'S3 (e)', shape='circle')
    dot.node('S4', 'S4 (i)', shape='circle')
    dot.node('S5', 'S5 (vowel + y)', shape='doublecircle') # Accepting state
    dot.node('S6', 'S6 (consonant)', shape='doublecircle') # Accepting state

    # Define transitions
    dot.edge('S1', 'S2', "y")
    dot.edge('S1', 'S3', "e")
    dot.edge('S2', 'S5', "vowel")
    dot.edge('S3', 'S4', "i")
    dot.edge('S4', 'S6', "consonant")

    return dot

# Test cases
test_words = ['boys', 'toys', 'ponies', 'skies', 'puppies', 'boies', 'toies', 'ponys',
'Babbies']
results = {word: is_plural_noun_accepted_fsa(word) for word in test_words}
print(results)

# Draw and render the FSA
fsa_graph = draw_fsa()
```

```
fsa_graph.render('plural_noun_fsa', format='png', cleanup=True)
```

```
from IPython.display import Image  
Image('plural_noun_fsa.png')
```

```
print(results)
```

SAMPLE INPUT-OUTPUT

```
{{'boys': True, 'toys': True, 'ponies': True, 'skies': True, 'puppies': True,  
  'boies': False, 'toies': False, 'ponys': False, 'Babbies': True}}
```




Figure 1: Finite State Automaton for Plural Noun Detection

Finite State Transducer

AIM

Design and implement a Finite State Transducer (FST) that accepts lexical forms of English words (e.g., shown below) and generates its corresponding plurals, based on the ****e-insertion spelling rule****:

$$\varepsilon \rightarrow e \quad / \quad \{x, s, z\}^{\wedge} _ s \#$$

\wedge is the morpheme boundary and $\#$ is the word boundary.

Input	Output
fox \wedge s#	foxes
boy \wedge s#	boys

PROGRAM

```
import graphviz

def generate_plural_fst(word):
    # Start in State 0
    state = 0

    # Transition logic based on states
    if state == 0:
        # Check for special cases requiring 'es'
        if word.endswith("^s#"):
            # Remove the suffix for processing
            base_word = word[:-3]

            # Check if the base word ends with special cases for pluralization
            if base_word.endswith(("x", "s", "z", "ch", "sh")):
                # Transition to plural form with 'es'
                return base_word + "es"
            else:
                # Default case: pluralize with 's'
                return base_word + "s"

    # If no rules apply, return the original word
    return word
```

```
def draw_fst():
    dot = graphviz.Digraph(comment='FST for Plural Generation')

    # Define states
    dot.node('S0', 'S0 (Start)', shape='circle') # Start state
    dot.node('S1', 'S1 (Special Case)', shape='circle')
    dot.node('S2', 'S2 (Default Case)', shape='circle')
    dot.node('Accept', 'Accepted', shape='doublecircle')

    # Define transitions
    dot.edge('S0', 'S1', 'ends with x/s/z')
    dot.edge('S1', 'Accept', 'append "es"')
    dot.edge('S0', 'S2', 'otherwise')
    dot.edge('S2', 'Accept', 'append "s"')

    return dot

# Test cases to check plural generation
test_cases = ["fox^s#", "boy^s#", "bus^s#", "dog^s#", "box^s#"]
# Dictionary to store results for each test case
plural_forms = {word: generate_plural_fst(word) for word in test_cases}

print(plural_forms)

# Draw and render the FST
fst_graph = draw_fst()
fst_graph.render('plural_generation_fst', format='png', cleanup=True)

# Display the generated FST diagram
from IPython.display import Image
Image('plural_generation_fst.png')
```

SAMPLE INPUT-OUTPUT

```
{'fox^s#': 'foxes', 'boy^s#': 'boys', 'bus^s#': 'buses', 'dog^s#': 'dogs', 'box^s#':  
  'boxes'}
```

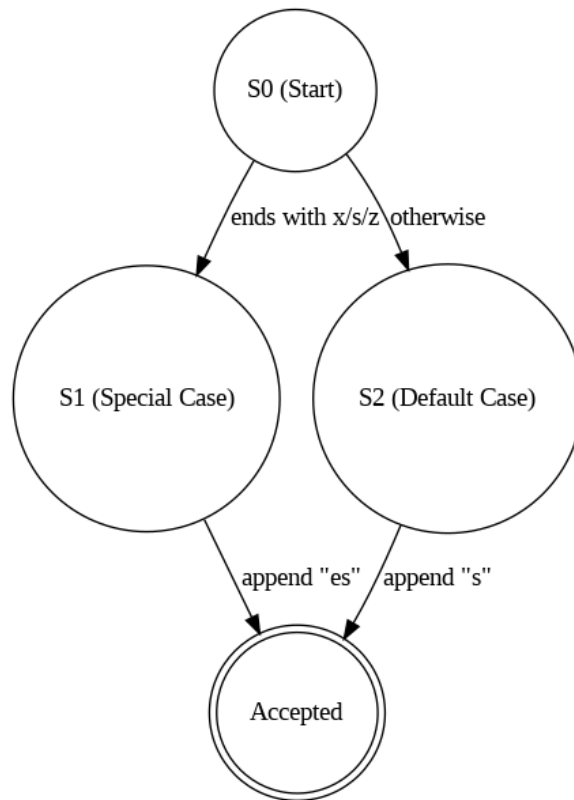


Figure 2: Finite State Transducer for Plural Noun Detection

Minimum Edit Distance

AIM

Implement the Minimum Edit Distance algorithm to find the edit distance between any two given strings. Also, list the edit operations.

PROGRAM

```
import numpy as np

def edit_distance(source, target):
    """
    Calculate the edit distance between source and target strings,
    and return the operations to convert source to target.
    """
    # Initialize dimensions
    len_source = len(source)
    len_target = len(target)

    # Initialize distance table
    distance_table = np.zeros((len_source + 1, len_target + 1), dtype=int)

    # Fill base cases
    for i in range(len_source + 1):
        distance_table[i][0] = i
    for j in range(len_target + 1):
        distance_table[0][j] = j

    # Fill distance table with dynamic programming
    for i in range(1, len_source + 1):
        for j in range(1, len_target + 1):
            if source[i - 1] == target[j - 1]:
                distance_table[i][j] = distance_table[i - 1][j - 1]
            else:
                distance_table[i][j] = min(
                    distance_table[i][j - 1] + 1,          # Insertion
                    distance_table[i - 1][j] + 1,          # Deletion
                    distance_table[i - 1][j - 1] + 2        # Substitution
                )

    # Trace back to determine operations
```

```
operations = []
i, j = len_source, len_target
while i > 0 or j > 0:
    if i > 0 and j > 0 and source[i - 1] == target[j - 1]:
        i, j = i - 1, j - 1
    elif i > 0 and distance_table[i][j] == distance_table[i - 1][j] + 1:
        operations.append(f"Delete '{source[i - 1]}' from position {i}")
        i -= 1
    elif j > 0 and distance_table[i][j] == distance_table[i][j - 1] + 1:
        operations.append(f"Insert '{target[j - 1]}' at position {j}")
        j -= 1
    else:
        operations.append(f"Substitute '{source[i - 1]}' with '{target[j - 1]}'
        at position {i}")
        i, j = i - 1, j - 1

# Reverse operations for correct order
operations.reverse()
return distance_table[len_source][len_target], operations

# Test the function with sample input
source = "exclusive"
target = "inclusive"
distance, operations = edit_distance(source, target)
print("Edit Distance:", distance)
print("Operations:")
for operation in operations:
    print(operation)
```

SAMPLE INPUT-OUTPUT

```
Edit Distance: 4
Operations:
Insert 'i' at position 1
Insert 'n' at position 2
Delete 'e' from position 1
Delete 'x' from position 2
```

Spell Checker

AIM

Design and implement a statistical spell checker for detecting and correcting non-word spelling errors in English, using the bigram language model. Your program should do the following:

1. Tokenize the corpus and create a vocabulary of unique words.
2. Create a bi-gram frequency table for all possible bigrams in the corpus.
3. Scan the given input text to identify the non-word spelling errors
4. Generate the candidate list using 1 edit distance from the misspelled words
5. Suggest the best candidate word by calculating the probability of the given sentence using the bigram LM.

PROGRAM

```
import re
from collections import defaultdict

def tokenize(text):
    # Tokenize text into words using regex for simplicity
    return re.findall(r'\b\w+\b', text.lower())

def build_vocabulary_and_bigrams(corpus):
    words = tokenize(corpus)
    vocabulary = set(words)

    bigram_freq = defaultdict(int)
    for i in range(len(words) - 1):
        bigram = (words[i], words[i+1])
        bigram_freq[bigram] += 1

    print("Vocabulary:", vocabulary) # Print the vocabulary
    print("Bigram Frequency Table:", dict(bigram_freq))
    # Print bigram frequency table

    return vocabulary, bigram_freq

def find_misspellings(text, vocabulary):
```

```
words = tokenize(text)
misspellings = [word for word in words if word not in vocabulary]
print("Identified Misspellings:", misspellings) # Print identified misspellings
return misspellings

def edits1(word):
    # Generate all words one edit away from 'word'
    letters = 'abcdefghijklmnopqrstuvwxyz'
    splits = [(word[:i], word[i:]) for i in range(len(word) + 1)]
    deletes = [L + R[1:] for L, R in splits if R]
    transposes = [L + R[1] + R[0] + R[2:] for L, R in splits if len(R) > 1]
    replaces = [L + c + R[1:] for L, R in splits if R for c in letters]
    inserts = [L + c + R for L, R in splits for c in letters]
    return set(deletes + transposes + replaces + inserts)

def get_candidates(word, vocabulary):
    # Get candidate corrections for the misspelled word
    candidates = {w for w in edits1(word) if w in vocabulary}
    print(f"Candidates for '{word}':", candidates) # Print candidate corrections
    return candidates

def get_multigram_candidates(misspelled_word, words, vocabulary):
    # Generate multi-word candidates by combining single-word candidates
    candidates = set()
    for i in range(len(words)):
        for j in range(i + 1, len(words) + 1):
            candidate = ' '.join(words[i:j])
            if candidate not in vocabulary:
                continue
            candidates.add(candidate)
    return candidates

def bigram_probability(sentence, bigram_freq):
    words = tokenize(sentence)
    bigram_prob = 1.0
    for i in range(len(words) - 1):
        bigram = (words[i], words[i+1])
        # Add a small constant to avoid zero probability
        bigram_prob *= (bigram_freq[bigram] + 1) / sum(bigram_freq.values())
    return bigram_prob
```

```
def correct_spelling(text, vocabulary, bigram_freq):
    words = tokenize(text)
    misspellings = find_misspellings(text, vocabulary)

    corrected_text = words[:]

    for misspelled_word in misspellings:
        candidates = get_candidates(misspelled_word, vocabulary)
        multi_candidates = get_multigram_candidates(misspelled_word, words, vocabulary)

        # Combine single-word and multi-word candidates
        all_candidates = candidates.union(multi_candidates)

        if all_candidates:
            best_candidate = max(
                all_candidates,
                key=lambda candidate: bigram_probability(
                    ' '.join(words).replace(misspelled_word, candidate), bigram_freq
                )
            )
            print(f"Replacing '{misspelled_word}' with '{best_candidate}'")
            # Print replacement info
            corrected_text = [best_candidate if w == misspelled_word else w for w
                              in corrected_text]

    return ' '.join(corrected_text)

# Example corpus and test case
corpus = (
    "Knowledge is power. The pen is mightier than the sword. Actions speak louder than words. "
    "Practice makes perfect. Better late than never. Birds of a feather flock together. "
    "A picture is worth a thousand words. When the going gets tough, the tough get going. "
    "Fortune favors the bold. Honesty is the best policy. Every cloud has a silver lining."
)

input_text = "Knwledge is powr. The pen is mightier than the sord. Practce makes prfect. A picure is worth a tousand words."
# Building vocabulary and bigram frequency table
vocabulary, bigram_freq = build_vocabulary_and_bigrams(corpus)
# Correct the input text
```



```
corrected_text = correct_spelling(input_text, vocabulary, bigram_freq)
print("Corrected Text:", corrected_text)
```

SAMPLE INPUT-OUTPUT

```
🔍 Vocabulary: {'policy', 'makes', 'has', 'gets', 'of', 'going', 'never', 'tough', 'pen', 'together', 'bold', 'favors', 'best', 'get',
Bigram Frequency Table: {('knowledge', 'is'): 1, ('is', 'power'): 1, ('power', 'the'): 1, ('the', 'pen'): 1, ('pen', 'is'): 1, ('is',
Identified Misspellings: ['knwledge', 'powr', 'mightier', 'sord', 'practce', 'prfect', 'picure', 'tousand']
Candidates for 'knwledge': {'knowledge'}
Replacing 'knwledge' with 'pen'
Candidates for 'powr': {'power'}
Replacing 'powr' with 'power'
Candidates for 'mightier': {'mightier'}
Replacing 'mightier' with 'mightier'
Candidates for 'sord': {'sword'}
Replacing 'sord' with 'pen'
Candidates for 'practce': {'practice'}
Replacing 'practce' with 'practice'
Candidates for 'prfect': {'perfect'}
Replacing 'prfect' with 'worth'
Candidates for 'picure': {'picture'}
Replacing 'picure' with 'picture'
Candidates for 'tousand': {'thousand'}
Replacing 'tousand' with 'thousand'
Corrected Text: pen is power the pen is mightier than the pen practice makes worth a picture is worth a thousand words
```

Sentiment Analysis

AIM

Implement a text classifier for sentiment analysis using the Naive Bayes theorem. Use Add-k smoothing to handle zero probabilities. Compare the performance of your classifier for k values 0.25, 0.75, and 1.

PROGRAM

```
import nltk
nltk.download('movie_reviews')
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import accuracy_score, classification_report
from nltk.corpus import movie_reviews

# Load positive and negative reviews
positive_fileids = movie_reviews.fileids('pos')
negative_fileids = movie_reviews.fileids('neg')

documents = [(movie_reviews.raw(fileid), 'pos') for fileid in positive_fileids]
documents.extend([(movie_reviews.raw(fileid), 'neg') for fileid in negative_fileids])

# Create DataFrame
df = pd.DataFrame(documents, columns=['text', 'label'])

# Split dataset into training and testing sets
X_train, X_test, y_train, y_test
= train_test_split(df['text'], df['label'], test_size=0.2, random_state=42)

# Vectorize text data
vectorizer = CountVectorizer()
X_train_vectorized = vectorizer.fit_transform(X_train)
X_test_vectorized = vectorizer.transform(X_test)

# Define Naive Bayes classifier
class NaiveBayes:
    def __init__(self, k=1):
        self.k = k
```

```
def fit(self, X, y):
    self.classes = np.unique(y)
    self.class_counts = {c: 0 for c in self.classes}
    self.feature_counts = {c: np.zeros(X.shape[1]) for c in self.classes}

    for i, c in enumerate(y):
        self.class_counts[c] += 1
        self.feature_counts[c] += X[i].toarray()[0]

    self.total_documents = len(y)
    self.vocab_size = X.shape[1]

def predict(self, X):
    predictions = []
    for i in range(X.shape[0]):
        post_probs = {}
        for c in self.classes:
            class_prob =
                (self.class_counts[c] + self.k) / (self.total_documents +
                self.k * len(self.classes))
            feature_prob = np.sum((X[i].toarray()[0] + self.k)
                / (self.feature_counts[c] + self.k * self.vocab_size))
            post_probs[c] = class_prob * feature_prob

        predictions.append(max(post_probs, key=post_probs.get))
    return np.array(predictions)

# Train and evaluate the model with different smoothing parameters
k_values = [0.25, 0.75, 1]
results = {}

for k in k_values:
    model = NaiveBayes(k=k)
    model.fit(X_train_vectorized, y_train)
    y_pred = model.predict(X_test_vectorized)

    accuracy = accuracy_score(y_test, y_pred)
    results[k] = accuracy
    print(f"Accuracy for k={k}: {accuracy:.4f}")
    print(classification_report(y_test, y_pred))
```

SAMPLE INPUT-OUTPUT

Accuracy for k=0.25: 0.4950					
	precision	recall	f1-score	support	
neg	0.00	0.00	0.00	201	
pos	0.50	0.99	0.66	199	
accuracy			0.49	400	
macro avg	0.25	0.50	0.33	400	
weighted avg	0.25	0.49	0.33	400	
Accuracy for k=0.75: 0.4975					
	precision	recall	f1-score	support	
neg	0.00	0.00	0.00	201	
pos	0.50	1.00	0.66	199	
accuracy			0.50	400	
macro avg	0.25	0.50	0.33	400	
weighted avg	0.25	0.50	0.33	400	

POS Tagging

AIM

Implement the Viterbi algorithm to find the most probable POS tag sequence for a given sentence, using the given probabilities:

a_{ij}	<i>STOP</i>	<i>NN</i>	<i>VB</i>	<i>JJ</i>	<i>RB</i>	b_{ik}	time	flies	fast
<i>START</i>	0	0.5	0.25	0.25	0	<i>NN</i>	0.1	0.01	0.01
<i>NN</i>	0.25	0.25	0.5	0	0	<i>VB</i>	0.01	0.1	0.01
<i>VB</i>	0.25	0.25	0	0	0.25	<i>JJ</i>	0	0	0.1
<i>JJ</i>	0	0.75	0	0.25	0	<i>RB</i>	0	0	0.1
<i>RB</i>	0.5	0.25	0	0.25	0				

PROGRAM

```
import numpy as np

def viterbi_algorithm(sentence, states, start_probabilities, transition_probs,
                       emission_probs):
    # Initialize the matrices
    num_states = len(states)
    num_words = len(sentence)

    viterbi = np.zeros((num_states, num_words))
    backtrack = np.zeros((num_states, num_words), dtype=int)

    # Initialization step
    for state_index, state in enumerate(states):
        viterbi[state_index, 0] = start_probabilities.get(state, 0)
        * emission_probs.get(state, {}).get(sentence[0], 0)

    # Recursion step
    for t in range(1, num_words):
        for current_state_index, current_state in enumerate(states):
            max_prob = 0
            prev_state_index = 0

            for previous_state_index, previous_state in enumerate(states):
                prob = (
                    viterbi[previous_state_index, t - 1]
                    * transition_probs.get(previous_state, {}).get(current_state, 0)
```

```
        * emission_probs.get(current_state, {}).get(sentence[t], 0)
    )
    if prob > max_prob:
        max_prob = prob
        prev_state_index = previous_state_index

    viterbi[current_state_index, t] = max_prob
    backtrack[current_state_index, t] = prev_state_index

# Termination step
final_prob = np.max(viterbi[:, -1])
final_state_index = np.argmax(viterbi[:, -1])

# Backtracking to find the best path
best_sequence = []
for t in range(num_words - 1, -1, -1):
    best_sequence.insert(0, states[final_state_index])
    final_state_index = backtrack[final_state_index, t]

return best_sequence, final_prob

# Example input data
states = ['NN', 'VB', 'JJ', 'RB']
sentence = ["time", "flies", "fast"]
start_probabilities = {'NN': 0.5, 'VB': 0.25, 'JJ': 0.25, 'RB': 0}
transition_probabilities = {
    'START': {'NN': 0.5, 'VB': 0.25, 'JJ': 0.25, 'RB': 0},
    'NN': {'NN': 0.25, 'VB': 0.5, 'JJ': 0, 'RB': 0},
    'VB': {'NN': 0.25, 'VB': 0.25, 'JJ': 0.25, 'RB': 0.25},
    'JJ': {'NN': 0.75, 'VB': 0, 'JJ': 0.25, 'RB': 0},
    'RB': {'NN': 0.25, 'VB': 0.25, 'JJ': 0, 'RB': 0.5},
}
emission_probabilities = {
    'NN': {'time': 0.1, 'flies': 0.01, 'fast': 0.01},
    'VB': {'time': 0.01, 'flies': 0.1, 'fast': 0.01},
    'JJ': {'time': 0, 'flies': 0, 'fast': 0.1},
    'RB': {'time': 0, 'flies': 0, 'fast': 0.1},
}

# Run the algorithm
best_sequence, best_probability = viterbi_algorithm(
```

```
        sentence,  
        states, start_probabilities, transition_probabilities, emission_probabilities  
    )
```

```
# Print results
```

```
print(f"Input Sentence: {' '.join(sentence)}")  
print(f"Best POS Tag Sequence: {' -> '.join(best_sequence)}")  
print(f"Probability of the Best Sequence: {best_probability:.10f}")
```

SAMPLE INPUT-OUTPUT

```
Input Sentence: time flies fast  
Best POS Tag Sequence: NN -> VB -> JJ  
Probability of the Best Sequence: 0.0000625000
```

Bigram Probability

AIM

Write a Python code to calculate bigrams from a given corpus and calculate the probability of any given sentence

PROGRAM

```
from collections import defaultdict
import math

def compute_bigrams(corpus):
    tokens = corpus.split() # Tokenize the text into words
    unigram_counts = defaultdict(int)
    bigram_counts = defaultdict(int)

    # Count unigrams and bigrams
    for idx in range(len(tokens)):
        unigram_counts[tokens[idx]] += 1
        if idx < len(tokens) - 1:
            bigram_counts[(tokens[idx], tokens[idx + 1])] += 1

    # Calculate bigram probabilities
    bigram_probs = {}
    for bigram, count in bigram_counts.items():
        bigram_probs[bigram] = count / unigram_counts[bigram[0]]

    return unigram_counts, bigram_counts, bigram_probs

def compute_sentence_probability(sentence, bigram_probs):
    words = sentence.split()
    probability = 1.0

    for idx in range(len(words) - 1):
        bigram = (words[idx], words[idx + 1])
        if bigram in bigram_probs:
            probability *= bigram_probs[bigram]
        else:
            probability *= 0.0001 # Assign small probability for unseen bigrams

    return probability
```



```
# Sample corpus and sentence
dataset = (
    "the quick brown fox jumps over the lazy dog "
    "a quick brown dog outpaces a quick fox"
)
input_text = "the quick brown fox jumps"

# Compute probabilities
unigrams, bigrams, probabilities = compute_bigrams(dataset)
text_probability = compute_sentence_probability(input_text, probabilities)

# Display results
print("Unigram Counts:", dict(unigrams))
print("\nBigram Counts:", dict(bigrams))
print("\nBigram Probabilities:", {k: round(v, 6) for k, v in probabilities.items()})
print(f"\nProbability of the sentence '{input_text}': {text_probability:.10f}")

}
```

SAMPLE INPUT-OUTPUT

Unigram Counts: {'the': 2, 'quick': 3, 'brown': 2, 'fox': 2, 'jumps': 1, 'over': 1, 'lazy': 1, 'dog': 2, 'a': 2, 'outpaces': 1}

Bigram Counts: {('the', 'quick'): 1, ('quick', 'brown'): 2, ('brown', 'fox'): 1, ('fox', 'jumps'): 1, ('jumps', 'over'): 1, ('over', 'the'): 1, ('the', 'lazy'): 1, ('lazy', 'dog'): 1, ('dog', 'a'): 1, ('a', 'quick'): 2, ('brown', 'dog'): 1, ('dog', 'outpaces'): 1, ('outpaces', 'a'): 1, ('quick', 'fox'): 1}

Bigram Probabilities: {('the', 'quick'): 0.5, ('quick', 'brown'): 0.666667, ('brown', 'fox'): 0.5, ('fox', 'jumps'): 0.5, ('jumps', 'over'): 1.0, ('over', 'the'): 1.0, ('the', 'lazy'): 0.5, ('lazy', 'dog'): 1.0, ('dog', 'a'): 0.5, ('a', 'quick'): 1.0, ('brown', 'dog'): 0.5, ('dog', 'outpaces'): 0.5, ('outpaces', 'a'): 1.0, ('quick', 'fox'): 0.333333}

Probability of the sentence 'the quick brown fox jumps': 0.0833333333

TF-IDF Matrix

AIM

Write a program to compute the TF-IDF matrix given a set of training documents. Also, calculate the cosine similarity between any two given documents or two given words.

PROGRAM

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
import numpy as np

def compute_tfidf_similarity(docs):
    """Calculate TF-IDF matrix and cosine similarities."""
    # Generate the TF-IDF matrix
    tfidf_vectorizer = TfidfVectorizer()
    tfidf_matrix = tfidf_vectorizer.fit_transform(docs)

    # Compute cosine similarity between first two documents
    cosine_sim = cosine_similarity(tfidf_matrix[0:1], tfidf_matrix[1:2])[0][0]

    return tfidf_vectorizer, tfidf_matrix, cosine_sim

def word_similarity(word1, word2, vectorizer, matrix):
    """Calculate cosine similarity between vectors of two words."""
    try:
        idx1, idx2 = vectorizer.vocabulary_[word1], vectorizer.vocabulary_[word2]
        word_vec1, word_vec2 = matrix[:, idx1].toarray().flatten(),
            matrix[:, idx2].toarray().flatten()
        return cosine_similarity([word_vec1], [word_vec2])[0][0]
    except KeyError:
        return 0.0 # Return zero if a word is not found in the vocabulary

# Example documents
sample_docs = [
    "time flies fast",
    "flies are flying fast",
    "fast time is good",
]

# Compute TF-IDF and similarities
```

```
vectorizer, tfidf_matrix, doc_similarity = compute_tfidf_similarity(sample_docs)
word_pair_similarity = word_similarity('time', 'flies', vectorizer, tfidf_matrix)

# Output
print("\nTF-IDF Matrix (as a Matrix):")
tfidf_matrix_array = tfidf_matrix.toarray()
terms = vectorizer.get_feature_names_out()

print("      " + " ".join(f"{term:10}" for term in terms))
for i, row in enumerate(tfidf_matrix_array):
    print(f"Doc {i+1}: " + " ".join(f"{val:.4f}" for val in row))

print("\nDocument Similarities:")
print(f"Similarity between Document 1 and Document 2: {doc_similarity:.4f}")

print("\nWord Similarities:")
print(f"Similarity between 'time' and 'flies': {word_pair_similarity:.4f}")
```

SAMPLE INPUT-OUTPUT

```
TF-IDF Matrix (as a Matrix):
      are      fast      flies      flying      good      is      time
Doc 1: 0.0000  0.4813  0.6198  0.0000  0.0000  0.0000  0.6198
Doc 2: 0.5845  0.3452  0.4445  0.5845  0.0000  0.0000  0.0000
Doc 3: 0.0000  0.3452  0.0000  0.0000  0.5845  0.5845  0.4445

Document Similarities:
Similarity between Document 1 and Document 2: 0.4417

Word Similarities:
Similarity between 'time' and 'flies': 0.6603
```

PPMI Matrix

AIM

Write a program to compute the PPMI matrix given a set of training documents. Also, calculate the cosine similarity between any two given documents or two given words.

PROGRAM

```
import numpy as np
from collections import defaultdict
from sklearn.metrics.pairwise import cosine_similarity

def generate_ppmi_matrix(corpus):
    """Generate PPMI matrix from a list of documents."""
    word_freq = defaultdict(int)
    co_occurrence = defaultdict(int)
    total_words = 0

    # Step 1: Calculate word frequencies and co-occurrence counts
    for sentence in corpus:
        words = sentence.split()
        total_words += len(words)
        for i, word in enumerate(words):
            word_freq[word] += 1
            for j in range(i + 1, len(words)):
                co_occurrence[(word, words[j])] += 1
                co_occurrence[(words[j], word)] += 1

    # Step 2: Build vocabulary and index
    vocabulary = sorted(word_freq.keys())
    vocab_size = len(vocabulary)
    word_to_index = {word: idx for idx, word in enumerate(vocabulary)}

    # Step 3: Create co-occurrence matrix
    co_occurrence_matrix = np.zeros((vocab_size, vocab_size))
    for (word1, word2), count in co_occurrence.items():
        idx1, idx2 = word_to_index[word1], word_to_index[word2]
        co_occurrence_matrix[idx1, idx2] = count

    # Step 4: Compute PPMI matrix
    ppmi_matrix = np.zeros_like(co_occurrence_matrix)
```

```
    for i in range(vocab_size):
        for j in range(vocab_size):
            joint_prob = co_occurrence_matrix[i, j] / total_words
            if joint_prob > 0:
                ppmi_value = np.log2(joint_prob / (word_freq[vocabulary[i]]
                                                    / total_words * word_freq[vocabulary[j]] / total_words))
                ppmi_matrix[i, j] = max(0, ppmi_value)

    return ppmi_matrix, word_to_index

def calculate_cosine_similarity(matrix, vocab_index, term1, term2):
    """Calculate cosine similarity between two words based on PPMI vectors."""
    idx1, idx2 = vocab_index[term1], vocab_index[term2]
    vector1, vector2 = matrix[idx1], matrix[idx2]
    return cosine_similarity([vector1], [vector2])[0][0]

# Training corpus (documents)
documents = [
    "the quick brown fox jumps over the lazy dog",
    "the quick brown fox is very quick",
    "the lazy dog sleeps all day",
]

# Step 1: Generate PPMI matrix and word-to-index mapping
ppmi_matrix, vocab_index = generate_ppmi_matrix(documents)

# Example 1: Cosine similarity between two words
word1 = "quick"
word2 = "lazy"
similarity_between_words = calculate_cosine_similarity(ppmi_matrix, vocab_index,
word1, word2)

# Step 2: Output the word similarity result
print(f"--- Word Similarity Analysis ---")
print(f"Words compared: '{word1}' and '{word2}'")
print(f"Cosine Similarity: {similarity_between_words:.4f}")
if similarity_between_words > 0.5:
    print("Interpretation: These words are quite similar in the context of the corpus.")
else:
    print("Interpretation: These words are not very similar in the context
of the corpus.")
```

```
# Example 2: Cosine similarity between two documents based on PPMI vectors
doc_ppmi_vectors = np.zeros((len(documents), ppmi_matrix.shape[0]))
for doc_idx, doc in enumerate(documents):
    words = doc.split()
    for word in words:
        if word in vocab_index:
            doc_ppmi_vectors[doc_idx] += ppmi_matrix[vocab_index[word]]

# Calculate document similarity
doc_similarity = cosine_similarity([doc_ppmi_vectors[0]], [doc_ppmi_vectors[1]])[0][0]

# Step 3: Output the document similarity result
print(f"\n--- Document Similarity Analysis ---")
print(f"Documents compared: Document 1 and Document 2")
print(f"Cosine Similarity: {doc_similarity:.4f}")
if doc_similarity > 0.7:
    print("Interpretation: These documents are highly similar in terms of
    content and context.")
elif doc_similarity > 0.4:
    print("Interpretation: These documents share some similarities but differ
    in context.")
else:
    print("Interpretation: These documents are quite different from each other in terms
    of content.")
```

SAMPLE INPUT-OUTPUT

```
--- Word Similarity Analysis ---
Words compared: 'quick' and 'lazy'
Cosine Similarity: 0.6016
Interpretation: These words are quite similar in the context of the corpus.

--- Document Similarity Analysis ---
Documents compared: Document 1 and Document 2
Cosine Similarity: 0.9336
Interpretation: These documents are highly similar in terms of content and context.
```

Naive Bayes Classifier

AIM

Implement a Naive Bayes classifier with add-1 smoothing using a given test data and disambiguate any word in a given test sentence. Use Bag-of-words as the feature. You may define your vocabulary. Sample Input :

No.	Sentence	Sense
1	I love fish. The smoked bass fish was delicious.	fish
2	The bass fish swam along the line.	fish
3	He hauled in a big catch of smoked bass fish.	fish
4	The bass guitar player played a smooth jazz line.	guitar

- Test Sentence: He loves jazz. The bass line provided the foundation for the guitar solo in the jazz piece
- Test word: bass
- Output: guitar

PROGRAM

```
import math
from collections import defaultdict

# Training data
training_examples = [
    {"text": "I love fish. The smoked bass fish was delicious.", "label": "fish"},
    {"text": "The bass fish swam along the line.", "label": "fish"},
    {"text": "He hauled in a big catch of smoked bass fish.", "label": "fish"},
    {"text": "The bass guitar player played a smooth jazz line.", "label": "guitar"},
]

# Test data
sentence_to_classify = "He loves jazz. The bass line provided the foundation for the guitar solo in the jazz piece"
target_word = "bass"

# 1. Tokenization and preparation of vocabulary and feature counts
def tokenize_text(text):
    return text.lower().replace(".", "").replace(",", "").split()
```

```
# Initialize vocabulary and word frequency counts for each sense
vocabulary = set()
sense_word_frequencies = defaultdict(lambda: defaultdict(int))
sense_labels = defaultdict(int)

# Process the training data
for example in training_examples:
    words = tokenize_text(example["text"])
    label = example["label"]
    sense_labels[label] += 1
    for word in words:
        if word != target_word: # Skip the ambiguous word
            vocabulary.add(word)
            sense_word_frequencies[label][word] += 1

# 2. Calculate probabilities using add-one smoothing (Laplace smoothing)
vocabulary_size = len(vocabulary)
total_documents = len(training_examples)
prior_probabilities = {label: count / total_documents for label, count in
sense_labels.items()}
word_likelihoods = defaultdict(lambda: defaultdict(float))

for label, word_counts in sense_word_frequencies.items():
    total_words_in_label = sum(word_counts.values())
    for word in vocabulary:
        word_likelihoods[label][word] = (word_counts[word] + 1)
        / (total_words_in_label + vocabulary_size)

# 3. Classify the test sentence
test_words = tokenize_text(sentence_to_classify)
log_posterior_probabilities = {}

for label in sense_labels.keys():
    # Start with the log of the prior probability
    log_posterior_probabilities[label] = math.log(prior_probabilities[label])
    for word in test_words:
        if word != target_word and word in vocabulary: # Skip the ambiguous word
            log_posterior_probabilities[label] += math.log(word_likelihoods[label][word])

# 4. Determine the most probable sense
predicted_label = max(log_posterior_probabilities, key=log_posterior_probabilities.get)
```



```
# Output with clear result and explanation
print(f"Predicted sense for the word '{target_word}' based on the test sentence:\n")
print(f"Test sentence: '{sentence_to_classify}'")
print(f"Predicted sense: {predicted_label}")
```

SAMPLE INPUT-OUTPUT

```
Predicted sense for the word 'bass' based on the test sentence:
Test sentence: 'He loves jazz. The bass line provided the foundation for the guitar solo in the jazz piece'
Predicted sense: guitar
```